

In [1]:

```
!pip install pytorch_tabnet
```

Collecting pytorch\_tabnet

Downloading pytorch\_tabnet-4.1.0-py3-none-any.whl.metadata (15 kB)

Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.11/dist-packages (from pytorch\_tabnet) (2.0.2)

Requirement already satisfied: scikit\_learn>0.21 in /usr/local/lib/python3.11/dist-packages (from pytorch\_tabnet) (1.6.1)

Requirement already satisfied: scipy>1.4 in /usr/local/lib/python3.11/dist-packages (from pytorch\_tabnet) (1.15.3)

Requirement already satisfied: torch>=1.3 in /usr/local/lib/python3.11/dist-packages (from pytorch\_tabnet) (2.6.0+cu124)

Requirement already satisfied: tqdm>=4.36 in /usr/local/lib/python3.11/dist-packages (from pytorch\_tabnet) (4.67.1)

Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.11/dist-packages (from scikit\_learn>0.21->pytorch\_tabnet) (1.5.1)

Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.11/dist-packages (from scikit\_learn>0.21->pytorch\_tabnet) (3.6.0)

Requirement already satisfied: filelock in /usr/local/lib/python3.11/dist-packages (from torch>=1.3->pytorch\_tabnet) (3.18.0)

Requirement already satisfied: typing-extensions>=4.10.0 in /usr/local/lib/python3.11/dist-packages (from torch>=1.3->pytorch\_tabnet) (4.14.0)

Requirement already satisfied: networkx in /usr/local/lib/python3.11/dist-packages (from torch>=1.3->pytorch\_tabnet) (3.5)

Requirement already satisfied: jinja2 in /usr/local/lib/python3.11/dist-packages (from torch>=1.3->pytorch\_tabnet) (3.1.6)

Requirement already satisfied: fsspec in /usr/local/lib/python3.11/dist-packages (from torch>=1.3->pytorch\_tabnet) (2025.3.2)

Collecting nvidia-cuda-nvrtc-cu12==12.4.127 (from torch>=1.3->pytorch\_tabnet)

Downloading nvidia\_cuda\_nvrtc\_cu12-12.4.127-py3-none-manylinux2014\_x86\_64.whl.metadata (1.5 kB)

Collecting nvidia-cuda-runtime-cu12==12.4.127 (from torch>=1.3->pytorch\_tabnet)

Downloading nvidia\_cuda\_runtime\_cu12-12.4.127-py3-none-manylinux2014\_x86\_64.whl.metadata (1.5 kB)

Collecting nvidia-cuda-cupti-cu12==12.4.127 (from torch>=1.3->pytorch\_tabnet)

Downloading nvidia\_cuda\_cupti\_cu12-12.4.127-py3-none-manylinux2014\_x86\_64.whl.metadata (1.6 kB)

Collecting nvidia-cudnn-cu12==9.1.0.70 (from torch>=1.3->pytorch\_tabnet)

Downloading nvidia\_cudnn\_cu12-9.1.0.70-py3-none-manylinux2014\_x86\_64.whl.metadata (1.6 kB)

Collecting nvidia-cublas-cu12==12.4.5.8 (from torch>=1.3->pytorch\_tabnet)

Downloading nvidia\_cublas\_cu12-12.4.5.8-py3-none-manylinux2014\_x86\_64.whl.metadata (1.5 kB)

Collecting nvidia-cufft-cu12==11.2.1.3 (from torch>=1.3->pytorch\_tabnet)

Downloading nvidia\_cufft\_cu12-11.2.1.3-py3-none-manylinux2014\_x86\_64.whl.metadata (1.5 kB)

Collecting nvidia-curand-cu12==10.3.5.147 (from torch>=1.3->pytorch\_tabnet)

Downloading nvidia\_curand\_cu12-10.3.5.147-py3-none-manylinux2014\_x86\_64.whl.metadata (1.5 kB)

Collecting nvidia-cusolver-cu12==11.6.1.9 (from torch>=1.3->pytorch\_tabnet)

Downloading nvidia\_cusolver\_cu12-11.6.1.9-py3-none-manylinux2014\_x86\_64.whl.metadata (1.6 kB)

Collecting nvidia-cusparselt-cu12==12.3.1.170 (from torch>=1.3->pytorch\_tabnet)

Downloading nvidia\_cusparselt\_cu12-12.3.1.170-py3-none-manylinux2014\_x86\_64.whl.metadata (1.6 kB)

Requirement already satisfied: nvidia-cusparse-cu12==12.3.1.170 in /usr/local/lib/python3.11/dist-packages (from torch>=1.3->pytorch\_tabnet) (12.3.1.170)

Requirement already satisfied: nvidia-nccl-cu12==2.21.5 in /usr/local/lib/python3.11/dist-packages (from torch>=1.3->pytorch\_tabnet) (2.21.5)

Requirement already satisfied: nvidia-nvtx-cu12==12.4.127 in /usr/local/lib/python3.11/dist-packages (from torch>=1.3->pytorch\_tabnet) (12.4.127)

Collecting nvidia-nvjitlink-cu12==12.4.127 (from torch>=1.3->pytorch\_tabnet)

Downloading nvidia\_nvjitlink\_cu12-12.4.127-py3-none-manylinux2014\_x86\_64.whl.metadata (1.5 kB)

Requirement already satisfied: triton==3.2.0 in /usr/local/lib/python3.11/dist-packages (from torch>=1.3->pytorch\_tabnet) (3.2.0)

Requirement already satisfied: sympy==1.13.1 in /usr/local/lib/python3.11/dist-packages (

```
from torch>=1.3->pytorch_tabnet) (1.13.1)
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.11/dist-packa
ges (from sympy==1.13.1->torch>=1.3->pytorch_tabnet) (1.3.0)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.11/dist-packages
(from jinja2->torch>=1.3->pytorch_tabnet) (3.0.2)
Downloading pytorch_tabnet-4.1.0-py3-none-any.whl (44 kB)
_____ 44.5/44.5 kB 2.1 MB/s eta 0:00:00
Downloading nvidia_cublas_cu12-12.4.5.8-py3-none-manylinux2014_x86_64.whl (363.4 MB)
_____ 363.4/363.4 MB 4.4 MB/s eta 0:00:00
Downloading nvidia_cuda_cupti_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl (13.8 MB)
_____ 13.8/13.8 MB 64.5 MB/s eta 0:00:00
Downloading nvidia_cuda_nvrtc_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl (24.6 MB)
_____ 24.6/24.6 MB 56.9 MB/s eta 0:00:00
Downloading nvidia_cuda_runtime_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl (883 kB)
_____ 883.7/883.7 kB 39.6 MB/s eta 0:00:00
Downloading nvidia_cudnn_cu12-9.1.0.70-py3-none-manylinux2014_x86_64.whl (664.8 MB)
_____ 664.8/664.8 MB 2.9 MB/s eta 0:00:00
Downloading nvidia_cufft_cu12-11.2.1.3-py3-none-manylinux2014_x86_64.whl (211.5 MB)
_____ 211.5/211.5 MB 6.6 MB/s eta 0:00:00
Downloading nvidia_curand_cu12-10.3.5.147-py3-none-manylinux2014_x86_64.whl (56.3 MB)
_____ 56.3/56.3 MB 13.3 MB/s eta 0:00:00
Downloading nvidia_cusolver_cu12-11.6.1.9-py3-none-manylinux2014_x86_64.whl (127.9 MB)
_____ 127.9/127.9 MB 8.5 MB/s eta 0:00:00
Downloading nvidia_cusparses_cu12-12.3.1.170-py3-none-manylinux2014_x86_64.whl (207.5 MB)
_____ 207.5/207.5 MB 6.3 MB/s eta 0:00:00
Downloading nvidia_nvjitlink_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl (21.1 MB)
_____ 21.1/21.1 MB 69.7 MB/s eta 0:00:00
Installing collected packages: nvidia-nvjitlink-cu12, nvidia-curand-cu12, nvidia-cufft-cu
12, nvidia-cuda-runtime-cu12, nvidia-cuda-nvrtc-cu12, nvidia-cuda-cupti-cu12, nvidia-cubl
as-cu12, nvidia-cusparses-cu12, nvidia-cudnn-cu12, nvidia-cusolver-cu12, pytorch_tabnet
Attempting uninstall: nvidia-nvjitlink-cu12
Found existing installation: nvidia-nvjitlink-cu12 12.5.82
Uninstalling nvidia-nvjitlink-cu12-12.5.82:
Successfully uninstalled nvidia-nvjitlink-cu12-12.5.82
Attempting uninstall: nvidia-curand-cu12
Found existing installation: nvidia-curand-cu12 10.3.6.82
Uninstalling nvidia-curand-cu12-10.3.6.82:
Successfully uninstalled nvidia-curand-cu12-10.3.6.82
Attempting uninstall: nvidia-cufft-cu12
Found existing installation: nvidia-cufft-cu12 11.2.3.61
Uninstalling nvidia-cufft-cu12-11.2.3.61:
Successfully uninstalled nvidia-cufft-cu12-11.2.3.61
Attempting uninstall: nvidia-cuda-runtime-cu12
Found existing installation: nvidia-cuda-runtime-cu12 12.5.82
Uninstalling nvidia-cuda-runtime-cu12-12.5.82:
Successfully uninstalled nvidia-cuda-runtime-cu12-12.5.82
Attempting uninstall: nvidia-cuda-nvrtc-cu12
Found existing installation: nvidia-cuda-nvrtc-cu12 12.5.82
Uninstalling nvidia-cuda-nvrtc-cu12-12.5.82:
Successfully uninstalled nvidia-cuda-nvrtc-cu12-12.5.82
Attempting uninstall: nvidia-cuda-cupti-cu12
Found existing installation: nvidia-cuda-cupti-cu12 12.5.82
Uninstalling nvidia-cuda-cupti-cu12-12.5.82:
Successfully uninstalled nvidia-cuda-cupti-cu12-12.5.82
Attempting uninstall: nvidia-cublas-cu12
Found existing installation: nvidia-cublas-cu12 12.5.3.2
Uninstalling nvidia-cublas-cu12-12.5.3.2:
Successfully uninstalled nvidia-cublas-cu12-12.5.3.2
Attempting uninstall: nvidia-cusparses-cu12
Found existing installation: nvidia-cusparses-cu12 12.5.1.3
Uninstalling nvidia-cusparses-cu12-12.5.1.3:
Successfully uninstalled nvidia-cusparses-cu12-12.5.1.3
Attempting uninstall: nvidia-cudnn-cu12
Found existing installation: nvidia-cudnn-cu12 9.3.0.75
Uninstalling nvidia-cudnn-cu12-9.3.0.75:
Successfully uninstalled nvidia-cudnn-cu12-9.3.0.75
Attempting uninstall: nvidia-cusolver-cu12
Found existing installation: nvidia-cusolver-cu12 11.6.3.83
Uninstalling nvidia-cusolver-cu12-11.6.3.83:
Successfully uninstalled nvidia-cusolver-cu12-11.6.3.83
Successfully installed nvidia-cublas-cu12-12.4.5.8 nvidia-cuda-cupti-cu12-12.4.127 nvidia
-cuda-nvrtc-cu12-12.4.127 nvidia-cuda-runtime-cu12-12.4.127 nvidia-cudnn-cu12-9.1.0.70 nv
```

idia-cufft-cu12-11.2.1.3 nvidia-curand-cu12-10.3.5.147 nvidia-cusolver-cu12-11.6.1.9 nvidia-cuspars-cu12-12.3.1.170 nvidia-nvjitlink-cu12-12.4.127 pytorch\_tabnet-4.1.0

In [2]:

```
import pandas as pd
import numpy as np
from sklearn.model_selection import StratifiedKFold, train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.metrics import accuracy_score, f1_score, precision_score, recall_score, roc_auc_score
from pytorch_tabnet.tab_model import TabNetClassifier
import torch
```

In [3]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.impute import SimpleImputer

df_uci = pd.read_csv("heart_disease_uci.csv")
df_uci.drop(columns=["id", "dataset"], inplace=True)
df_uci["target"] = df_uci["num"].apply(lambda x: 1 if x > 0 else 0)
df_uci.drop(columns=["num"], inplace=True)

numerical_uci = ['age', 'trestbps', 'chol', 'thalch', 'oldpeak', 'ca']
categorical_uci = ['sex', 'cp', 'fbs', 'restecg', 'exang', 'slope', 'thal']
```

In [4]:

```
df_uci[numerical_uci] = SimpleImputer(strategy="median").fit_transform(df_uci[numerical_uci])
for col in categorical_uci:
    df_uci[col] = LabelEncoder().fit_transform(df_uci[col].astype(str))
df_uci[numerical_uci] = StandardScaler().fit_transform(df_uci[numerical_uci])

# Renaming columns
df_uci.rename(columns={"chol": "cholesterol", "sex": "gender"}, inplace=True)
```

In [5]:

```
# Preprocessing
df_kaggle = pd.read_csv("cardio_train.csv", sep=';')
df_kaggle.drop(columns=["id"], inplace=True)
df_kaggle["target"] = df_kaggle["cardio"]
df_kaggle.drop(columns=["cardio"], inplace=True)

categorical_kaggle = ['gender', 'cholesterol', 'gluc', 'smoke', 'alco', 'active']
numerical_kaggle = ['age', 'height', 'weight', 'ap_hi', 'ap_lo']

for col in categorical_kaggle:
    df_kaggle[col] = LabelEncoder().fit_transform(df_kaggle[col].astype(str))
df_kaggle[numerical_kaggle] = StandardScaler().fit_transform(df_kaggle[numerical_kaggle])
```

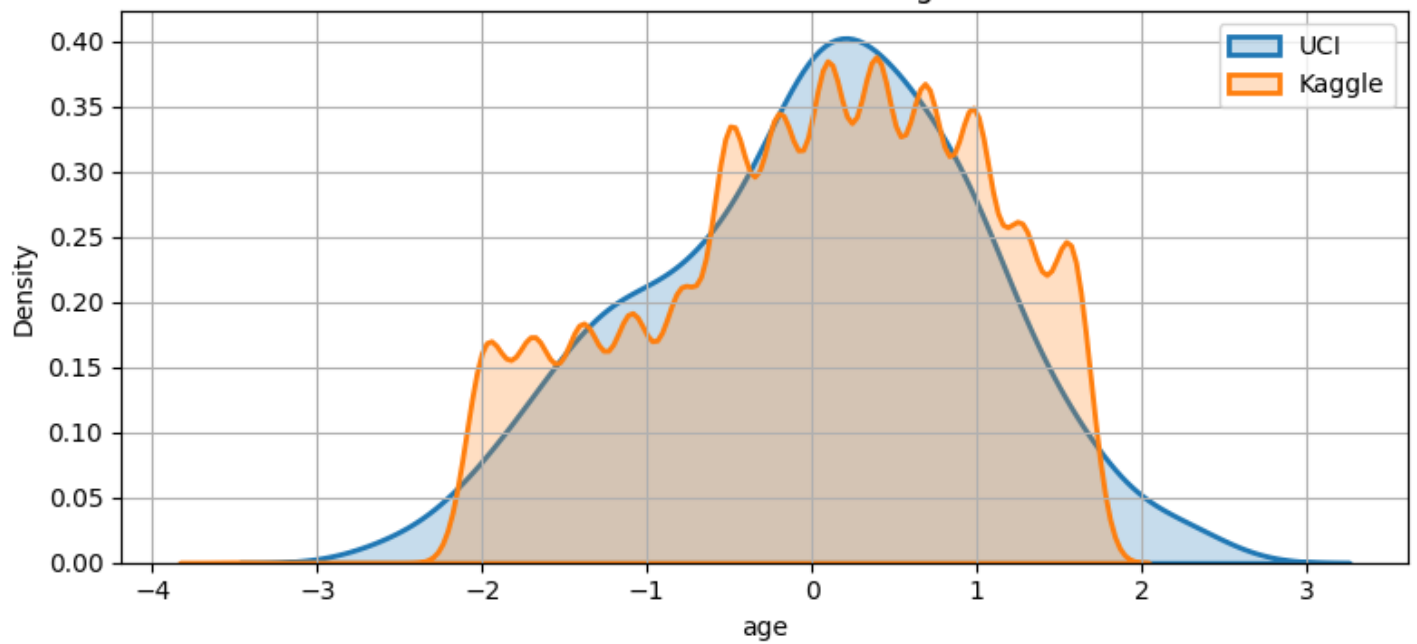
In [6]:

```
# Plotting
overlapping_features = ['age', 'cholesterol', 'gender']

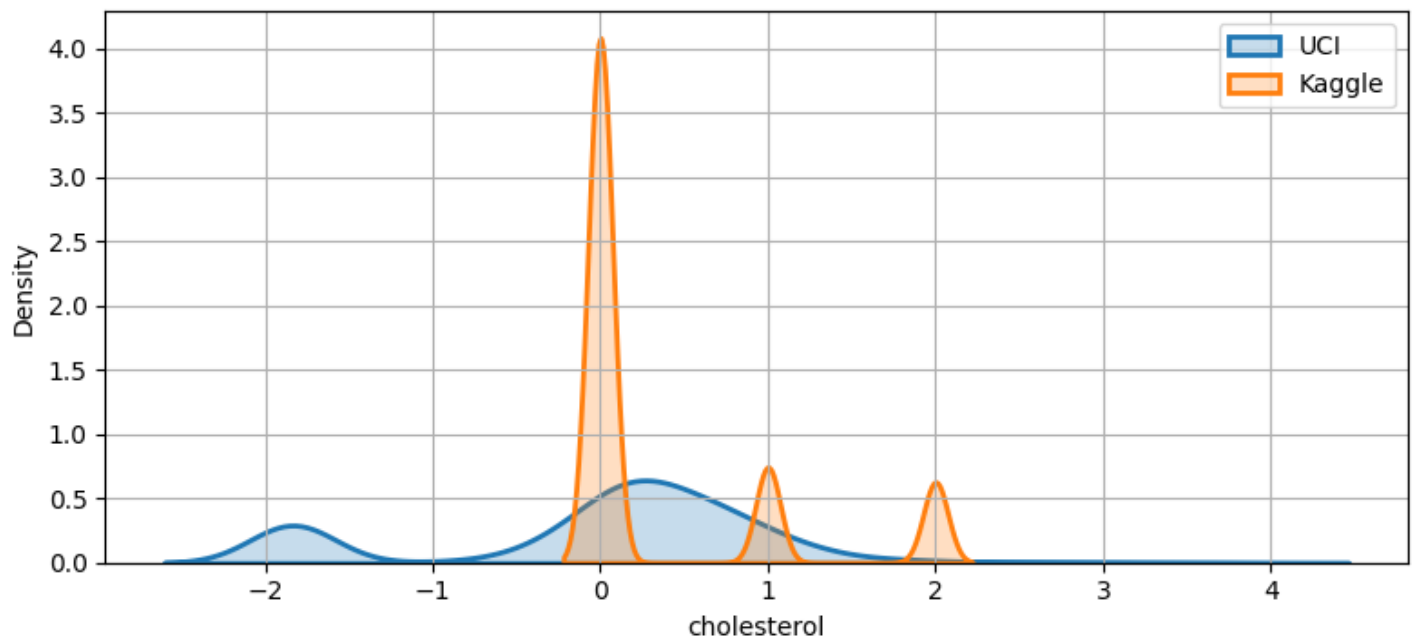
for feature in overlapping_features:
    plt.figure(figsize=(8, 4))
    sns.kdeplot(df_uci[feature], label='UCI', fill=True, linewidth=2)
    sns.kdeplot(df_kaggle[feature], label='Kaggle', fill=True, linewidth=2)
    plt.title(f"Feature Distribution: {feature}")
    plt.xlabel(feature)
    plt.legend()
    plt.grid(True)
```

```
plt.tight_layout()
plt.show()
```


Feature Distribution: age



Feature Distribution: cholesterol



Feature Distribution: gender



In [7]:

```
from pytorch_tabnet.tab_model import TabNetClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, roc_auc_score
import torch

# Data
X_uci = df_uci.drop(columns=["target"]).values
y_uci = df_uci["target"].values

X_uci_train, X_uci_test, y_uci_train, y_uci_test = train_test_split(
    X_uci, y_uci, test_size=0.2, stratify=y_uci, random_state=42
)
```

In [8]:

```
# Training on UCI
tabnet_uci = TabNetClassifier(
    n_d=32, n_a=32, n_steps=5,
    gamma=1.5, lambda_sparse=1e-3, momentum=0.5,
    optimizer_fn=torch.optim.Adam,
    optimizer_params=dict(lr=0.01),
    scheduler_params={"step_size": 10, "gamma": 0.9},
    scheduler_fn=torch.optim.lr_scheduler.StepLR,
    verbose=0,
    seed=42
)

tabnet_uci.fit(
    X_uci_train, y_uci_train,
    eval_set=[(X_uci_train, y_uci_train), (X_uci_test, y_uci_test)],
    eval_name=['train', 'val'],
    eval_metric=['accuracy'],
    max_epochs=100,
    patience=10,
    batch_size=1024,
    virtual_batch_size=128,
    num_workers=0,
    drop_last=False
)
```

Early stopping occurred at epoch 30 with best\_epoch = 20 and best\_val\_accuracy = 0.7663

/usr/local/lib/python3.11/dist-packages/pytorch\_tabnet/callbacks.py:172: UserWarning: Best weights from best epoch are automatically used!  
warnings.warn(wrn\_msg)

In [9]:

```
# Training Kaggle
X_kaggle_data = df_kaggle.drop(columns=["target"]).values
y_kaggle_data = df_kaggle["target"].values

X_kaggle_train, X_kaggle_test, y_kaggle_train, y_kaggle_test = train_test_split(
    X_kaggle_data, y_kaggle_data, stratify=y_kaggle_data, test_size=0.2, random_state=42
)

tabnet_kaggle = TabNetClassifier(
    n_d=32, n_a=32, n_steps=5,
    gamma=1.5, lambda_sparse=1e-3, momentum=0.5,
    optimizer_fn=torch.optim.Adam,
    optimizer_params=dict(lr=0.01),
    scheduler_params={"step_size": 10, "gamma": 0.9},
    scheduler_fn=torch.optim.lr_scheduler.StepLR,
    verbose=0,
    seed=42
)
```

```

tabnet_kaggle.fit(
    X_kaggle_train, y_kaggle_train,
    eval_set=[(X_kaggle_train, y_kaggle_train), (X_kaggle_test, y_kaggle_test)],
    eval_name=['train', 'val'],
    eval_metric=['accuracy'],
    max_epochs=100,
    patience=10,
    batch_size=1024,
    virtual_batch_size=128,
    num_workers=0,
    drop_last=False
)

```

Early stopping occurred at epoch 19 with best\_epoch = 9 and best\_val\_accuracy = 0.71736

```

/usr/local/lib/python3.11/dist-packages/pytorch_tabnet/callbacks.py:172: UserWarning: Best weights from best epoch are automatically used!
  warnings.warn(wrn_msg)

```

In [10]:

```

import shap
X_sample_uci = X_uci_test[:200]
background_uci = X_uci_train[:100]
explainer_uci = shap.KernelExplainer(lambda x: tabnet_uci.predict_proba(x)[:, 1], background_uci)
shap_values_uci = explainer_uci.shap_values(X_sample_uci)

```

In [11]:

```

X_sample_kaggle = X_kaggle_test[:200]
background_kaggle = X_kaggle_train[:100]
explainer_kaggle = shap.KernelExplainer(lambda x: tabnet_kaggle.predict_proba(x)[:, 1], background_kaggle)
shap_values_kaggle = explainer_kaggle.shap_values(X_sample_kaggle)

```

In [12]:

```

feature_names_uci = df_uci.drop(columns=["target"]).columns.tolist()
feature_names_kaggle = df_kaggle.drop(columns=["target"]).columns.tolist()

mean_abs_shap_uci = np.abs(shap_values_uci).mean(axis=0)
mean_abs_shap_kaggle = np.abs(shap_values_kaggle).mean(axis=0)

```

In [13]:


```

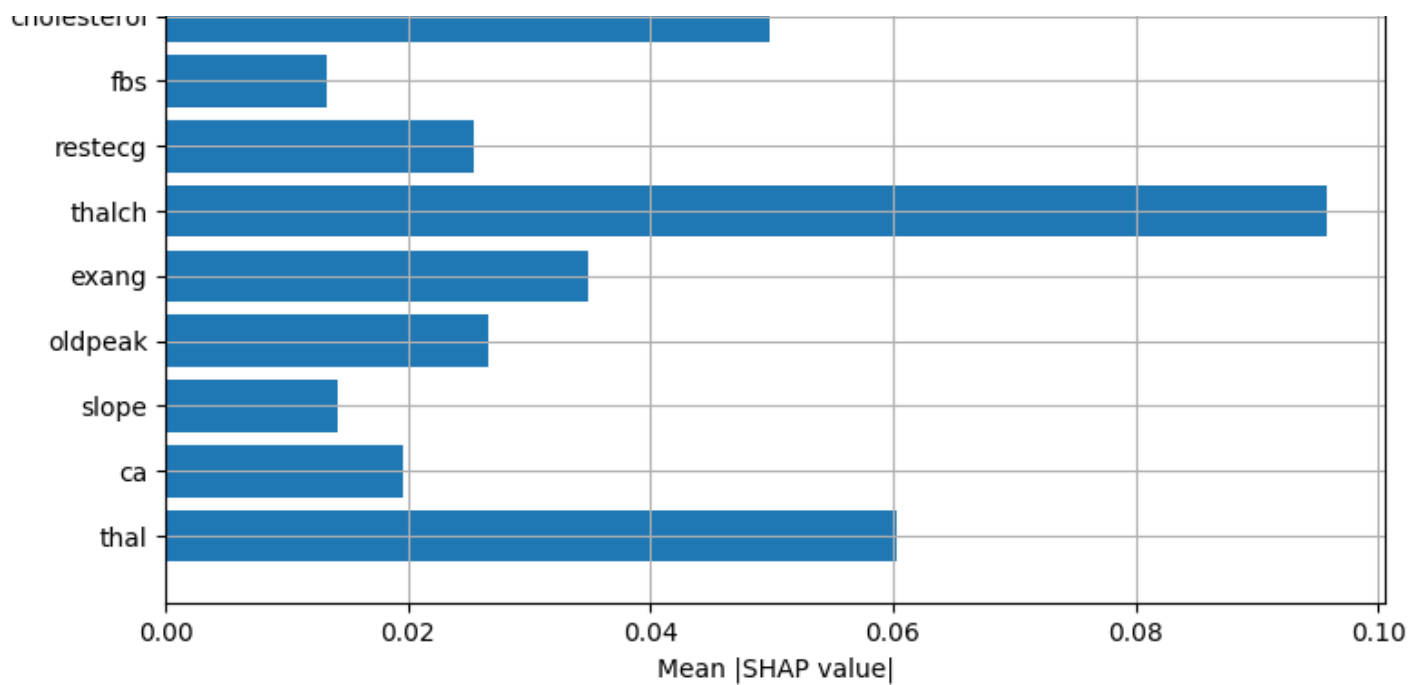
import matplotlib.pyplot as plt
import numpy as np

# Plotting UCI SHAP
plt.figure(figsize=(8, 6))
plt.barh(feature_names_uci[::-1], mean_abs_shap_uci[::-1])
plt.title("UCI-trained TabNet SHAP Importance")
plt.xlabel("Mean |SHAP value|")
plt.grid(True)
plt.tight_layout()
plt.show()

```

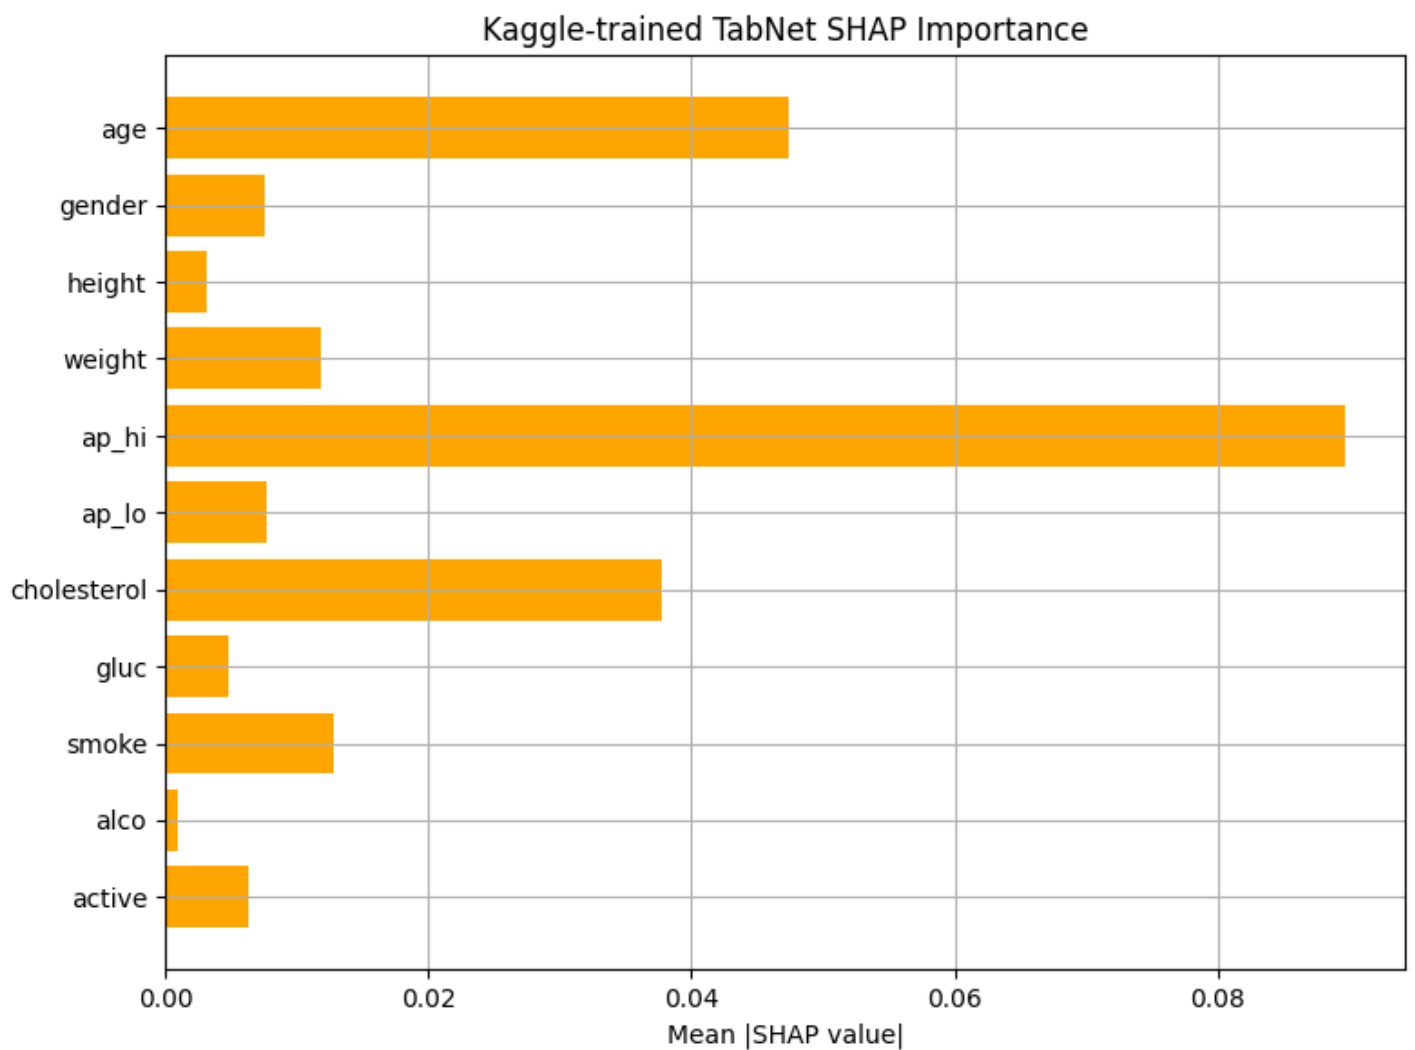
UCI-trained TabNet SHAP Importance





In [14]:

```
# Plotting kaggle SHAP
plt.figure(figsize=(8, 6))
plt.barh(feature_names_kaggle[::-1], mean_abs_shap_kaggle[::-1], color='orange')
plt.title("Kaggle-trained TabNet SHAP Importance")
plt.xlabel("Mean |SHAP value|")
plt.grid(True)
plt.tight_layout()
plt.show()
```



In [ ]:

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

uci_names = df_uci.drop(columns=["target"]).columns.tolist()
kaggle_names = df_kaggle.drop(columns=["target"]).columns.tolist()

shared_features = list(set(uci_names).intersection(set(kaggle_names)))

# Mapping SHAP values to names
shap_uci_series = pd.Series(mean_abs_shap_uci, index=uci_names)
shap_kaggle_series = pd.Series(mean_abs_shap_kaggle, index=kaggle_names)

# Keeping only shared features
shap_comparison = pd.DataFrame({
    'UCI': shap_uci_series[shared_features],
    'Kaggle': shap_kaggle_series[shared_features]
})

# 5. Sorting
shap_comparison = shap_comparison.loc[shap_comparison.mean(axis=1).sort_values(ascending=True).index]

# 6. Plot grouped horizontal bar chart
shap_comparison.plot(kind='barh', figsize=(10, 6))
plt.title("SHAP Importance: UCI vs Kaggle (Shared Features)")
plt.xlabel("Mean |SHAP value|")
plt.grid(True)
plt.tight_layout()
plt.show()

```

In [15]:

```

# Real feature names
uci_names = df_uci.drop(columns=["target"]).columns.tolist()
kaggle_names = df_kaggle.drop(columns=["target"]).columns.tolist()

all_features = sorted(set(uci_names).union(set(kaggle_names)))

# Mapping SHAP values
shap_uci_series = pd.Series(mean_abs_shap_uci, index=uci_names)
shap_kaggle_series = pd.Series(mean_abs_shap_kaggle, index=kaggle_names)

# Reindexing
shap_uci_series_full = shap_uci_series.reindex(all_features, fill_value=0)
shap_kaggle_series_full = shap_kaggle_series.reindex(all_features, fill_value=0)

# Combining into DataFrame
shap_all_df = pd.DataFrame({
    'UCI': shap_uci_series_full,
    'Kaggle': shap_kaggle_series_full
})

# Sorting by total SHAP value
shap_all_df = shap_all_df.loc[shap_all_df.sum(axis=1).sort_values(ascending=True).index]

shap_all_df.plot(kind='barh', figsize=(8, 6)) # smaller size
plt.title("SHAP Importance Comparison", fontsize=12)
plt.xlabel("Mean |SHAP value|", fontsize=10)
plt.xticks(fontsize=8)
plt.yticks(fontsize=8)
plt.tight_layout()
plt.savefig("shap_compact.png", dpi=300, bbox_inches='tight')
plt.show()


```

SHAP Importance Comparison









In [ ]: